

In [80]: `print("Hello World...!!!")`

Hello World...!!!

In [81]: `import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from mlxtend.plotting import plot_confusion_matrix
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline`

In [82]: `df = pd.read_csv("Social_Network_Ads.csv")`

In [83]: `df`

Out[83]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

Basic Stats

In [84]: `df.head()`

Out[84]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

In [85]: `df.describe()`

Out[85]:

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

In [86]: `df.shape`

Out[86]: (400, 5)


In [87]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         400 non-null   int64
1   Gender          400 non-null   object
2   Age             400 non-null   int64
3   EstimatedSalary 400 non-null   int64
4   Purchased       400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```


```
In [88]: df.isna().sum()
```

```
Out[88]: User ID      0
Gender      0
Age         0
EstimatedSalary  0
Purchased   0
dtype: int64
```

```
In [89]: histplot = sns.histplot(df['Age'], kde=True, bins=30, color='red', alpha=0.3)
for i in histplot.containers:
    histplot.bar_label(i)
plt.show()
```



```
In [90]: histplot = sns.histplot(df['EstimatedSalary'], kde=True, bins=30, color='red',
for i in histplot.containers:
    histplot.bar_label(i)
plt.show()
```



```
In [91]: df["Gender"].value_counts()
```

```
Out[91]: Female    204  
Male        196  
Name: Gender, dtype: int64
```


```
In [92]: def gender_encoder(value):  
        if (value == "Male"):  
            return 1  
        elif (value == "Female"):  
            return 0  
        else:  
            return -1
```

```
In [93]: df["Gender"] = df["Gender"].apply(gender_encoder)
```


```
In [94]: df["Purchased"].value_counts()
```

```
Out[94]: 0    257  
        1    143  
        Name: Purchased, dtype: int64
```


```
In [95]: countplot = sns.countplot(x=df["Purchased"])  
        for i in countplot.containers:  
            countplot.bar_label(i)  
        plt.show()
```



```
In [96]: countplot = sns.countplot(x=df["Purchased"], hue=df["Gender"], palette="twilig
for i in countplot.containers:
    countplot.bar_label(i,)
plt.show()
```



```
In [99]: sns.heatmap(df.corr(), annot=True)
plt.show()
```



Data Preparation

```
In [104]: x = df[["Age", "EstimatedSalary"]]
y = df["Purchased"]
```

```
In [105]: scaler = StandardScaler()
x = scaler.fit_transform(x)
```

```
In [106]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [107]: x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
Out[107]: ((320, 2), (80, 2), (320,), (80,))
```

Model Building

```
In [34]: model = LogisticRegression(n_jobs=-1)
```

```
In [35]: model.fit(x_train, y_train)
```

```
Out[35]: LogisticRegression(n_jobs=-1)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.


```
In [37]: y_pred = model.predict(x_test)
```

Evaluation

```
In [38]: cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[50  2]
 [ 9 19]]
```

```
In [39]: plot_confusion_matrix(conf_mat=cm, figsize=(5,5), show_normed=True)
plt.show()
```



```
In [40]:
```

```
print(f"TN value is {cm[0][0]}")
print(f"FP value is {cm[0][1]}")
print(f"FN value is {cm[1][0]}")
print(f"TP value is {cm[1][1]}")
```

```
TN value is 50
FP value is 2
FN value is 9
TP value is 19
```

```
In [47]: print(f"Accuracy score is {accuracy_score(y_test, y_pred)}")
```

```
Accuracy score is 0.8625
```

```
In [46]: print(f"Error rate is {1-accuracy_score(y_test, y_pred)}")
```

```
Error rate is 0.13749999999999996
```

```
In [43]: print(f"Precision score is {precision_score(y_test, y_pred)}")
```

```
Precision score is 0.9047619047619048
```

```
In [44]: print(f"Recall score is {recall_score(y_test, y_pred)}")
```

```
Recall score is 0.6785714285714286
```

In [45]:

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.96	0.90	52
1	0.90	0.68	0.78	28
accuracy			0.86	80
macro avg	0.88	0.82	0.84	80
weighted avg	0.87	0.86	0.86	80

In []: